

# Artificial Bee Colony (ABC) Algorithm for Regression

## COMP6000 Project

Bradley Fuller

[bf84@kent.ac.uk](mailto:bf84@kent.ac.uk)



School of Computing

University of Kent

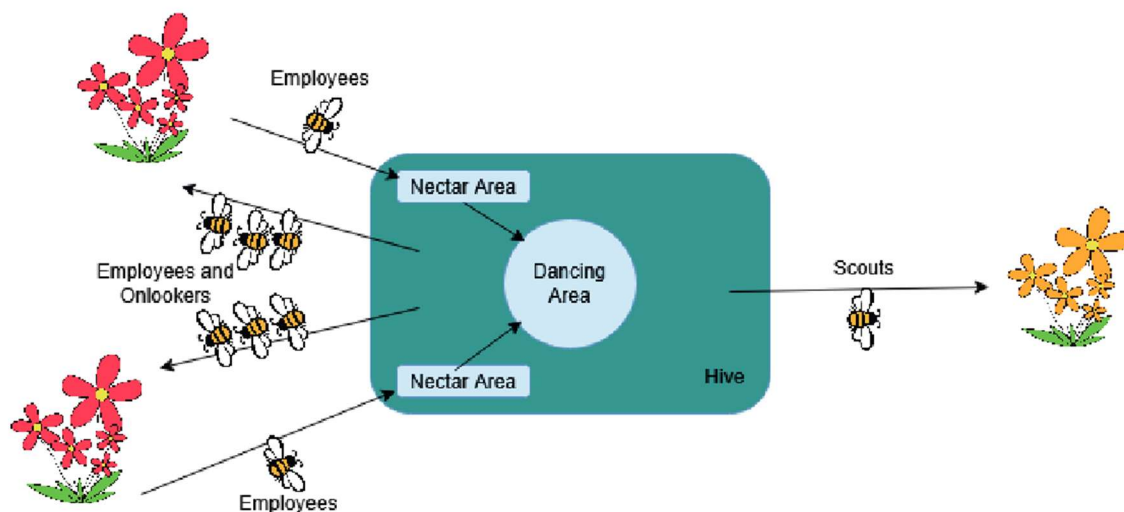
22<sup>nd</sup> March 2025

## 2 1 Abstract

3 This report will go into detail about the concept, development and uses of the Artificial Bee  
4 Colony algorithm, especially, as the title suggests, its use in solving regression problems. The  
5 algorithm itself is inspired by the foraging behaviour of honeybees, such as their exploration  
6 and exploitation strategies. Due to the nature of what it is based on, it has the benefit of  
7 being understandable by those with knowledge in this field, and those who do not. It will  
8 also prove its use and effectiveness in function optimisation, with direct examples of its use.  
9 The results showing clearly its efficiency in solving the function optimisation problem and  
10 the regression problem (a real world application as shown in this report), with the speed of  
11 which the correct or optimal solution is found. In this project I have tried the linear  
12 regression model but I am aware of other regression models that could be explored in the  
13 future, such as the exponential model.

## 2 Introduction

As stated earlier, the artificial bee colony algorithm is an optimisation algorithm that is based upon honeybee foraging behaviour. The 'bees' are assigned three different roles which in the case of the program, align to different tasks being performed, with food sources being different possible results, as illustrated below. To give a brief explanation for each role, the employee bee takes a known source and searches nearby, meaning to slightly alter its current source and see if it results in a higher fitness, if yes that one is saved, if not it is discarded and it sticks with its original source. The goal of this phase is to broaden the search and hopefully find better results. The onlooker bee phase is more focused, with each bee picking a source at random, with the chance being weighted towards ones of a higher fitness, before just performing the role of employee. The goal of this phase is similar to employee, however with the chance leaning more towards those of higher weighting, it makes sure the better quality sources get explored more, improving them. Finally is the scout phase, which unlike the first two which run for every iteration, only runs after a set requirement is met, that being when a source is 'exhausted'. In the code, when a source is checked for better nearby ones, and nothing better is found, a token on that source slowly goes up, and once it hits the limit, the source is exhausted, and is likely close to, if not the best, in that area. As such, the scout generates an entirely new source in its place, and the employees then start the search again. Obviously the best overall source would be saved on every iteration, which by the end would be returned as the result. The frequency at which these tasks are performed combined with the amount of times it is iterated upon allows it to find optimal solutions very quickly, as will be shown more in the Function Optimisation and Regression sections. On top of this, we also simulate it three times and show the average of the results, due to the level of randomness that comes with the algorithm. To elaborate on the randomness aspect, when the program starts, the food sources are generated randomly within certain boundaries. As such, sometimes a source could be generated that is already very close to the solution, meaning technically it didn't take many iterations, or it could start far from it and take a little bit before finding the optimal solution.



### 3.1 Function Optimisation

To start with, the objective of function optimisation is to find the values in a function which provide the best result, or the highest 'fitness'. In the case of the functions we tested, the goal in each case was to minimise the objective function, meaning that we wanted to take the inverse of the value, and get the smallest result possible, as the smaller the value the closer it would be to the solution to the function, as such after calculating the fitness for each function, it would be inverted and compared as if it was smaller instead. Now, obviously this projects focus is on regression, however focusing on function optimisation first was beneficial to us when starting off. The main reason being that when using a function to test the algorithm, unlike the regression data we used, we already knew what the result should be. As such, we could tell when the programming was running well as once it was returning results that lined up with the actual solution for the function, we were able to tell the program was working correctly. We started with a simple function, as once the basis is written, changing the code to work for other functions only required editing the function calculation aspect, the algorithm itself would be unchanged.

Once the starting function was being returned correctly, while this in of itself was a good sign, we decided to test further by using multiple 'optimisation test problems'. We took multiple problems from here and implemented them into the function, even keeping all of them and swapping between with a simple switch case, to allow for showing the results of multiple functions later. From here, moving on to regression wasn't too huge of a step, as the ABC algorithm itself was working, though the algorithm function did need very minor changes, the bigger focus was the fitness function, which would need the most changes to work with the regression model we were working with, this will be delved into more in the Regression section. With the results found from the function optimisation we could at least tell how efficient the program could be, with there not needing to be a huge amount of iterations to get close to if not the perfect result, even when taking into account the levels of randomness that comes with the algorithm. In terms of the specific functions chosen, the ones we ended up using were the Beale, Booth, Matyas and Three Hump Camel. These ones were each chosen for their own reasons, the first two were chosen as they were still somewhat simple and not too far from the starting function we used for testing, which was  $f(x,y) = (x + y)^2 + (x - y)^2$ . The second two being a bit tough to implement and typically taking slightly longer to figure out the optimal solution, albeit not by much however.

## 3.2 Function Optimisation | Results

Below is the results from running the Beale Function. Obviously to show every single iteration would take up too much space, but the results at least show how close the program would get to the optimal solution with not many iterations.

```
Simulation 1...
Simulation 1 Done! Performed 100 iterations.
Simulation 2...
Simulation 2 Done! Performed 100 iterations.
Simulation 3...
Simulation 3 Done! Performed 100 iterations.
Results of Beale Function:
Simulations      Best Solution (x,y) found      Computed f(x,y) of best solution
1                (2.9054, 0.4799)              0.0020
2                (2.8904, 0.4668)              0.0026
3                (3.0176, 0.5053)              0.0001
| | | | |      Average Performance              0.0016

[Done] exited with code=0 in 1.049 seconds
```

In the case of the Beale Function, the correct solution is (3, 0.5), so the program is getting close, which is where the computed of best solution part helps to tell. This is figured out by dividing one by the best fitness, and subtracting that by one. This changes the result from a fitness score to instead represent roughly how far off the value is from the best solution, as well as taking the average from each simulation. Below is the same result table for the 'Three Hump Camel' function. The answer to which is (0,0)

```
Simulation 1...
Simulation 1 Done! Performed 100 iterations.
Simulation 2...
Simulation 2 Done! Performed 100 iterations.
Simulation 3...
Simulation 3 Done! Performed 100 iterations.
Results of Three Hump Camel Function:
Simulations      Best Solution (x,y) found      Computed f(x,y) of best solution
1                (0.0000, -0.0000)              0.0000
2                (0.0000, 0.0000)              0.0000
3                (0.0000, -0.0000)              0.0000
| | | | |      Average Performance              0.0000

[Done] exited with code=0 in 1.393 seconds
```

As stated earlier, while not even a second difference, the three hump camel function does take very slightly longer to iterate over the other function.

## 89 4.1 Regression | Data Set

90 Once the algorithm was functioning correctly for function optimisation, the next part was  
91 taking that and using it to find the optimal solution for a regression problem. The data set I  
92 used was the Real Estate valuation data from  
93 <https://archive.ics.uci.edu/dataset/477/real+estate+valuation+data+set>. This data set has six  
94 features (transaction date, house age, distance to nearest mrt, no. of convenience stores,  
95 latitude, longitude) which is used to predict the house price (target variable). Notice that the  
96 range of different features varies dramatically, for example the distance to the nearest MRT  
97 and the number of convenience stores. Therefore, after importing the data file into the  
98 program, I normalise the data set. By normalising, I scale all of the data to ensure its range is  
99 between 0 and 1. This was done by going through every column of data, and finding the  
100 biggest and smallest in each one. These values are set as the 'min' and 'max' for each, and  
101 then normalise the data in each column according to the formula below:

102 
$$\text{'normalisedData'} = (\text{data} - \text{min}) / (\text{max} - \text{min})$$

103 The data set has 414 examples, in order to get more accurate prediction results, I put 2/3 of  
104 the examples into the training set, and the remaining 1/3 into testing set. The training data is  
105 the data that will be input into the ABC algorithm to find optimal weights in the linear  
106 regression model, to predict the house price. The testing data is used to evaluate the  
107 predictive accuracy of the linear regression model on the house price. While talking about  
108 the data set, I thought I would mention the simulation aspect. When the program runs, the  
109 ABC Algorithm is performed 3 times fully, and averages are taken, to help account for the  
110 randomness factor, however in the case for the data set being normalised and randomly  
111 split, this is performed once and used 3 times for the data. I felt it unnecessary to also  
112 randomise the data assigning as there is already such a large randomness factor at play with  
113 the algorithm itself. This is something I could look into with cross validation which will be  
114 touched on more in the Conclusion.

## 115 4.2 Regression | Applying ABC

116 With the data normalised and split, I can now use it within the algorithm itself. As was stated  
117 in the function optimisation section earlier, the ABC Algorithm uses 'food sources' to search  
118 around and identify potential solutions, in the case of regression, these food sources are the  
119 weights for the linear regression model, the formula for which is shown below.



$$120 \quad y(x_1, x_2, x_3, x_4, x_5, x_6) = w_0 + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + w_5x_5 + w_6x_6$$

121 In this case, the x's in the function are the values from the training data set for each column,  
122 and as stated the w's (or weights) are the food source from ABC. The vast majority of the  
123 ABC algorithm itself plays out much the same as how it does with function optimisation,  
124 however I did have to make a few changes for it to work correctly on this function. Changes  
125 such as changing food sources to have 6 values instead of 2, increasing the amount of  
126 iterations to allow for more accurate results, and making sure the data is read from  
127 specifically the training set. The part I had to redo completely, was the fitness function, since  
128 I was now working with a completely different function entirely. This function especially,  
129 required working with something completely different, while for every function optimisation  
130 problem I could have the numbers scored based on their size, in that case since they were all  
131 minimisation, it gave better scores to lower numbers, the regression problem is different.

132 I want to measure how close a predicted value was to the actual value, so while fitness could  
133 work here, it was better to work with Mean Squared Error (MSE for short). MSE is useful for  
134 finding the difference between two values, especially in that the values getting squared  
135 makes it more punishing for larger errors. In the case of this function, the actual value and  
136 predicted values are the ones I want to calculate the difference between. This is done by  
137 filling in the function with the values being called in the function, so filling in the x's with the  
138 features from the data set, the w's being the food sources (predicted values) and the y being  
139 the actual value (the cost). Then once we have the two values, just taking the predicted  
140 value away from the actual value and squaring the result, which results in the MSE. Now  
141 when this is used in ABC, it still requires the number to be larger as higher value is what the  
142 'bees' search for, as such we invert this number with a simple function,  $(1/(1+mse))$ , which  
143 first makes sure the number is between 0 and 1, and helps the 'bees' see whether this is an  
144 optimal source. Once all the iterations are done, this same function is performed with the  
145 test data and the average MSE is compared to see if the results match up well.

## 4.3 Regression | Results

The results of the regression function lined up with what I was hoping to achieve with this which I was very happy to see. It is however a little bit more nuanced than just showing the result table. As such, below I have plotted two line charts, the first one titled; “Best MSE” is somewhat self-explanatory. I took what was considered the best MSE in each iteration from one simulation and plotted it. In this case it was for 200 iterations, which is higher than function optimisation to allow for a more detailed answer, but not as high as it could go as after a certain point it feels pointless as the line does not move much. The second, and more interesting chart is the Average MSE chart, this one, like the best MSE chart, takes from the same simulation and same number of iterations, but instead takes the average MSE of every food source at that point in the iteration. At first glance they will seem quite similar, however looking at the Y axis on both, the average MSE starts from a much higher value, in this case close to 7, whereas Best starts from around 0.9, by the end however, both are close to 0. I feel this well demonstrates both the randomness factor mentioned earlier, with how high the starting point for the average MSE is as many of the early sources will be extremely inefficient, but also shows how efficient ABC can be at finding the optimal solutions, as even towards the end the average is somewhat close to the best MSE, thus, the optimal solution.



To go alongside these charts, below is the result table of the same function. The two charts specifically being based on the data from ‘Simulation 1’.

Results:			
Simulations	Best Source Found	Training MSE	Test MSE
1	-0.0354, 0.0166, -0.1817, 0.1933, 0.4225, 0.1346, 0.1706	0.0107	0.0157
2	0.5403, -0.0564, -0.3000, -0.5651, -0.1275, -0.2616, 0.2505	0.0170	0.0244
3	-0.2097, 0.0914, -0.0391, 0.0653, 0.1172, 0.0670, 0.5653	0.0085	0.0139
Averages:		0.0121	0.0180

[Done] exited with code=0 in 2.625 seconds

From this result table, we can confirm that even with only 200 iterations, the algorithm is finding solutions that very accurately predict what the price of the house could be. While the best sources found are nice to see, it in of itself doesn’t tell us much, but the Training and Test MSE give insight into how close that prediction is. For example with Simulation 1 again, the training MSE prediction is 0.0107 off of the actual value, and that is after being squared to enhance any errors. This clearly shows an accurate prediction, which for the training data is maybe a bit expected, which is why I have the testing data. The Test MSE is only 0.0157 off, again after being squared. While a bit further off than the Training MSE, it is still within decimals of being correct, and this is using sources found via the training data, which proves its use in predicting.



180 As an extra to include, while I feel the above examples illustrate the efficiency well, below I  
 181 have included two extra images. Both of these are after upping the number of iterations  
 182 from 200 to 1000 per simulation, obviously letting the algorithm search for much longer. The  
 183 first is just a selection of iterations from simulation 1 in this case which I wanted to include  
 184 to show the jump for average MSE a bit more accurately. As for the second, it is the results  
 185 of these much longer iterations. While it isn't the perfect judge of efficiency I wanted to  
 186 include the time the code took to run, as an added way of showing its speed.

Iteration Number: 1	Best MSE: 1.2177142896477235	Average MSE: 6.455341328614682
Iteration Number: 2	Best MSE: 1.124431343543768	Average MSE: 5.438445963241946
Iteration Number: 3	Best MSE: 0.9236754093681538	Average MSE: 4.870251801089357
Iteration Number: 4	Best MSE: 0.9146326453395164	Average MSE: 4.399231028811036
Iteration Number: 5	Best MSE: 0.7703737316766992	Average MSE: 3.848917696993566
Iteration Number: 6	Best MSE: 0.7703737316766992	Average MSE: 3.4038982277111067
Iteration Number: 7	Best MSE: 0.7703737316766992	Average MSE: 3.024032244926464
Iteration Number: 8	Best MSE: 0.7703737316766992	Average MSE: 2.7793311083394676
Iteration Number: 9	Best MSE: 0.7703737316766992	Average MSE: 2.642999474528046
Iteration Number: 10	Best MSE: 0.7627847472559033	Average MSE: 2.539291576369453
Iteration Number: 11	Best MSE: 0.5429328755680323	Average MSE: 2.362575176114073
Iteration Number: 12	Best MSE: 0.4720620465473009	Average MSE: 2.2477990432139237
Iteration Number: 13	Best MSE: 0.4720620465473009	Average MSE: 2.153019292744582
Iteration Number: 14	Best MSE: 0.4645723013929035	Average MSE: 2.038393124542688
Iteration Number: 15	Best MSE: 0.4293475439443539	Average MSE: 1.9603967518099985
Iteration Number: 16	Best MSE: 0.4283564586164019	Average MSE: 1.897767860751884
Iteration Number: 17	Best MSE: 0.4283564586164019	Average MSE: 1.8109894452985182
Iteration Number: 18	Best MSE: 0.4283564586164019	Average MSE: 1.7483861903165177
Iteration Number: 19	Best MSE: 0.26885056924049744	Average MSE: 1.6851481194005422
Iteration Number: 20	Best MSE: 0.26885056924049744	Average MSE: 1.632490824709182
Iteration Number: 21	Best MSE: 0.26885056924049744	Average MSE: 1.5601567571219621
Iteration Number: 22	Best MSE: 0.26885056924049744	Average MSE: 1.5314832654868176
Iteration Number: 23	Best MSE: 0.26736095227442025	Average MSE: 1.4821733775024581
Iteration Number: 24	Best MSE: 0.26736095227442025	Average MSE: 1.4656007238493514
Iteration Number: 25	Best MSE: 0.26026473780212545	Average MSE: 1.4404886145381122

Results:			
Simulations	Best Source Found	Training MSE	Test MSE
1	0.1659, 0.0515, -0.1072, -0.2512, 0.0855, 0.3707, -0.0681	0.0075	0.0065
2	0.3018, 0.0211, -0.1206, -0.3546, 0.0558, 0.1501, -0.0244	0.0070	0.0059
3	0.5543, 0.0360, -0.1044, -0.5672, 0.1228, -0.0949, -0.2601	0.0083	0.0070
	Averages:	0.0076	0.0065
[Done] exited with code=0 in 5.213 seconds			

## 187 5 Conclusion

188 In conclusion, I feel I have succeeded in proving the effectiveness in using the ABC algorithm  
189 to solve both function optimisation and regression problems. I believe the speed at which  
190 the algorithm converges onto the optimal solutions in both are strong evidence of its  
191 efficiency. As well as this, the regression model having such low MSE values, especially on  
192 the runs with more iterations, shows that its predictions are extremely accurate, even when  
193 compared to data it is not trained on at all and even taking into account the inherent level of  
194 randomness that comes with using it. I believe its efficiency is one of the major strengths of  
195 this algorithm. I would also like to touch on its versatility, it's something I mentioned earlier  
196 on in this report, as when the program was first written, its main focus was on the function  
197 optimisation aspect, however when I moved onto using it for regression, while the fitness  
198 function did change pretty much entirely, the algorithm itself had very minor changes. I feel  
199 this is a huge advantage of using this, as I think it wouldn't take too many changes to make it  
200 work for other functions.

201 While I feel the algorithm is efficient and all things considered not very expensive in terms of  
202 runtime, I do have to factor that while it adding more iterations did improve results, it does  
203 come at the cost of runtime. With the data and functions I was working with this wouldn't  
204 be an issue, however I can see that if working with larger data sets or just adding more  
205 iterations and simulations, this could start to add up, not even factoring working on more  
206 complex models such as the exponential model. While with this data I do not believe it  
207 made much of, if not any, difference, the randomness is something that for larger data sets  
208 could be more of an issue, and in that situation would need to be looked into. If I were to  
209 continue working on this in the future, trying larger data sets and different models and  
210 seeing if ABC would work as effectively with those as it did for these ones. As mentioned in  
211 an earlier section, I would also like to see if cross validation would have any impact on the  
212 result. In this case that would mean instead of using the same sorted data for all 3  
213 simulations, instead, sorting them differently every time, maybe even changing the  
214 distribution of examples such as 50/50, or trying with less training and more testing. I feel  
215 having those as different simulations then comparing those results would give a good insight  
216 to its impact. Less related to ABC, if I came back to this in the future, I would be interested to  
217 try different algorithms, such as Particle Swarm and Ant Colony, as I wasn't able to look into  
218 those as much as I would've liked for comparison due to time constraints.

219 Overall, I do feel that the ABC algorithm was efficient in this use case, and did a good job of  
220 solving the function optimisation and regression problems, especially with how close the  
221 predictions in regression were.

## 222 6 Acknowledgements | References

223 I would like to thank Jie Shao, my project supervisor, for her assistance and patience during  
224 this project, with the advice given being a huge factor in my understanding of this concept  
225 early on, as well as the Natural Computation lectures from herself and Alex Freitas, the  
226 lecture slides I frequently used as a refresher. As well I would like to thank the references  
227 attached below, some of which supplied data and others used for insight and research.

228 Surjanovic, S. and Bingham, D. (2013) *Virtual library of simulation experiments: Optimization*  
229 *Test Functions and Datasets*. Available at: <https://www.sfu.ca/~ssurjano/optimization.html>  
230 (Accessed: 18 February 2025).

231 Where I obtained the optimisation test functions.

232 Yeh, I.-C. (2018) *Real estate valuation, UCI Machine Learning Repository*. Available at:  
233 <https://archive.ics.uci.edu/dataset/477/real+estate+valuation+data+set> (Accessed: 16  
234 February 2025).

235 Where I obtained the data set used for regression.